

# Consensus Protocol

We decided to implement the RAFT protocol for this assignment. We followed the directions described and set up 8 nodes running as go routines. Each node is modelled as a state machine so that it can move between follower, candidate or leader states depending on what action it just took.

Our algorithm works pretty well and elects a leader within 0.5ms in most cases. Below we have a screenshot of nodes going up for candidacy, receiving majority votes, leaving leadership and another leader being promptly elected 3 times.

```
$ ./RAFT.exe
ID: 7 up for candidacy
Node 0 voting for node 7
Node 1 voting for node 7
Node 6 voting for node 7
Node 3 voting for node 7
Node 2 voting for node 7
Node 4 voting for node 7
Node 5 voting for node 7
ID: 7 elected as leader
ID: 7 dropping leadership
ID: 6 up for candidacy
Node 3 voting for node 6
ID: 4 up for candidacy
Node 5 voting for node 6
Node 0 voting for node 6
Node 7 voting for node 6
ID: 2 up for candidacy
Node 1 voting for node 6
Node 2 voting for node 6
ID: 6 elected as leader
Node 4 voting for node 6
ID: 6 dropping leadership
Node 6 voting for node 4
ID: 0 up for candidacy
Node 1 voting for node 0
Node 4 voting for node 0
Node 2 voting for node 0
Node 3 voting for node 0
Node 7 voting for node 0
Node 5 voting for node 0
Node 6 voting for node 0
ID: 0 elected as leader
```

Figure 1: Example of our RAFT program running with print outputs.

Occasionally, there will be a lot of rounds of elections as multiple nodes try to become leaders all at once but our program resolves itself and elects a leader at the end and will never have multiple leaders.

The code can be built using:

Go build RAFT.go

Then run the executable produced.